

# Lecture 3: Markov Decision Processes

Bolei Zhou

The Chinese University of Hong Kong

*bzhou@ie.cuhk.edu.hk*

January 14, 2020

# This week's Plan


## ① Last Time

- ① Key elements of an RL agent: model, value, policy

## ② This Week: Decision Making in MDP

- ① Markov decision processes (MDP)
- ② Policy evaluation in MDP
- ③ Control in MDP: policy iteration and value iteration

# Markov Decision Process (MDP)



- ① Markov Decision Process can model a lot of real-world problem. It formally describes the framework of reinforcement learning
- ② Under MDP, the environment is fully observable.
  - ① Optimal control primarily deals with continuous MDPs
  - ② Partially observable problems can be converted into MDPs

# Define the model of the environment

- Markov Processes
- Markov Reward Processes(MRPs)
- Markov Decision Processes (MDPs)

# Markov Property

- ① The history of states:  $h_t = \{s_1, s_2, s_3, \dots, s_t\}$
- ② State  $s_t$  is Markov if and only if:

$$p(s_{t+1}|s_t) = p(s_{t+1}|h_t) \quad (1)$$

$$p(s_{t+1}|s_t, a_t) = p(s_{t+1}|h_t, a_t) \quad (2)$$

- ③ "The future is independent of the past given the present"


# Markov Process/Markov Chain

- ①  $P$  is the dynamics/transition model that specifies  $p(s_{t+1} = s' | s_t = s)$
- ② State transition matrix: (From|To)

$$P = \begin{bmatrix} P(s_1|s_1) & P(s_2|s_1) & \dots & P(s_N|s_1) \\ P(s_1|s_2) & P(s_2|s_2) & \dots & P(s_N|s_2) \\ \vdots & \vdots & \ddots & \vdots \\ P(s_1|s_N) & P(s_2|s_N) & \dots & P(s_N|s_N) \end{bmatrix}$$

- ③ Note that there are no rewards or no actions.

# Example of MP



① Samples episodes starting from  $s_3$


- ①  $s_3, s_4, s_5, s_6, s_6$
- ②  $s_3, s_2, s_3, s_2, s_1$
- ③  $s_3, s_4, s_4, s_5, s_5$

# Markov Reward Process (MRP)

- ① Markov Reward Process is a Markov Chain + reward
- ② Definition of Markov Reward Process (MRP)
  - ①  $S$  is a (finite) set of states ( $s \in S$ )
  - ②  $P$  is dynamics/transition model that specifies  $P(S_{t+1} = s' | s_t = s)$
  - ③  $R$  is a reward function  $R(s_t = s) = \mathbb{E}[r_t | s_t = s]$
  - ④ Discount factor  $\gamma \in [0, 1]$
- ③ If finite number of states,  $R$  can be a vector



## Example of MRP



Reward:  $+5$  in  $s_1$ ,  $+10$  in  $s_7$ ,  $0$  in all other states. So that we can represent  $R = [5, 0, 0, 0, 0, 0, 10]$

# Return function and Value function

## 1 Definition of Horizon

- 1 Number of maximum time steps in each episode
- 2 Can be infinite, otherwise called finite Markov (reward) Process

## 2 Definition of Return

- 1 Discounted sum of rewards from time step  $t$  to horizon

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots + \gamma^{T-t-1} R_T$$

## 3 Definition of state value function $V_t(s)$ for a MRP

- 1 Expected return from  $t$  in state  $s$


$$\begin{aligned} V_t(s) &= \mathbb{E}[G_t | s_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T | s_t = s] \end{aligned}$$

- 2 Present value of future rewards

# Why Discount Factor $\gamma$

- ① Mathematically convenient to discount rewards
- ② Avoids infinite returns in cyclic Markov processes
- ③ Uncertainty about the future may not be fully represented
- ④ If the reward is financial, immediate rewards may earn more interest than delayed rewards
- ⑤ Animal/human behaviour shows preference for immediate reward
- ⑥ It is sometimes possible to use undiscounted Markov reward processes (i.e.  $\gamma = 1$ ), e.g. if all sequences terminate.
  - ①  $\gamma = 0$ : Only care about the immediate reward
  - ②  $\gamma = 1$ : Future reward is equal to the immediate reward.

# Example of MRP



- ① Reward: +5 in  $s_1$ , +10 in  $s_7$ , 0 in all other states. So that we can represent  $R = [5, 0, 0, 0, 0, 0, 10]$
- ② Sample returns for a 4-step episodes with  $\gamma = 1/2$ 
  - ① return for  $s_4, s_5, s_6, s_7$  :  $0 + \frac{1}{2} \times 0 + \frac{1}{4} \times 0 + \frac{1}{8} \times 10 = 1.25$
  - ② return for  $s_4, s_3, s_2, s_1$  :  $0 + \frac{1}{2} \times 0 + \frac{1}{4} \times 0 + \frac{1}{8} \times 5 = 0.625$
  - ③ return  $s_4, s_5, s_6, s_7$  : = 0
  - ④ How to compute the value function?

# Computing the Value of a Markov Reward Process

- 1 Value function: expected return from starting in state  $s$

$$V(s) = \mathbb{E}[G_t | s_t = s] = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | s_t = s]$$

- 2 MRP value function satisfies the following Bellman equation:

$$V(s) = \underbrace{R(s)}_{\text{Immediate reward}} + \underbrace{\gamma \sum_{s' \in S} P(s'|s) V(s')}_{\text{Discounted sum of future reward}}$$

- 3 Practice: To derive the Bellman equation for  $V(s)$

# Matrix Form of Bellman Equation for MRP

Therefore, we can express  $V(s)$  using the matrix form:

$$\begin{bmatrix} V(s_1) \\ V(s_2) \\ \vdots \\ V(s_N) \end{bmatrix} = \begin{bmatrix} R(s_1) \\ R(s_2) \\ \vdots \\ R(s_N) \end{bmatrix} + \gamma \begin{bmatrix} P(s_1|s_1) & P(s_2|s_1) & \dots & P(s_N|s_1) \\ P(s_1|s_2) & P(s_2|s_2) & \dots & P(s_N|s_2) \\ \vdots & \vdots & \ddots & \vdots \\ P(s_1|s_N) & P(s_2|s_N) & \dots & P(s_N|s_N) \end{bmatrix} \begin{bmatrix} V(s_1) \\ V(s_2) \\ \vdots \\ V(s_N) \end{bmatrix}$$

$$V = R + \gamma PV$$

① Analytic solution for value of MRP:  $V = (I - \gamma P)^{-1}R$

① Matrix inverse takes the complexity  $O(N^3)$  for  $N$  states

② Only possible for a small MRPs

# Iterative Algorithm for Computing Value of a MRP

- ① Iterative methods for large MRPs:
  - ① Dynamic Programming
  - ② Monte-Carlo evaluation
  - ③ Temporal-Difference learning

# Iterative Algorithm for Computing Value of a MRP

---

**Algorithm 1** Monte Carlo simulation to calculate MRP value function

---

- 1:  $i \leftarrow 0, G_t \leftarrow 0$
  - 2: **while**  $i \neq N$  **do**
  - 3:   generate an episode, starting from state  $s$  and time  $t$
  - 4:   Using the generated episode, calculate return  $g = \sum_{i=t}^{H-1} \gamma^{i-t} r_i$
  - 5:    $G_t \leftarrow G_t + g, i \leftarrow i + 1$
  - 6: **end while**
  - 7:  $V_t(s) \leftarrow G_t / N$
-



# Iterative Algorithm for Computing Value of a MRP

---

**Algorithm 2** Iterative algorithm to calculate MRP value function

---

- 1: for all states  $s \in S$ ,  $V'(s) \leftarrow 0$ ,  $V(s) \leftarrow \infty$
  - 2: **while**  $\|V - V'\| > \epsilon$  **do**
  - 3:    $V \leftarrow V'$
  - 4:   For all states  $s \in S$ ,  $V'(s) = R(s) + \gamma \sum_{s' \in S} P(s'|s)V(s')$
  - 5: **end while**
  - 6: return  $V'(s)$  for all  $s \in S$
-

# Markov Decision Process (MDP)

- ① Markov Decision Process is Markov Reward Process with decisions.
- ② Definition of MDP
  - ①  $S$  is a finite set of states
  - ②  $A$  is a finite set of actions
  - ③  $P^a$  is dynamics/transition model for each action
$$P(s_{t+1} = s' | s_t = s, a_t = a)$$
  - ④  $R$  is a reward function  $R(s_t = s, a_t = a) = \mathbb{E}[r_t | s_t = s, a_t = a]$
  - ⑤ Discount factor  $\gamma \in [0, 1]$
- ③ MDP is a tuple:  $(S, A, P, R, \gamma)$

# Policy in MDP



- ① Policy specifies what action to take in each state
- ② Give a state, specify a distribution over actions
- ③ Policy:  $\pi(a|s) = P(a_t = a | s_t = s)$
- ④ Policies are stationary (time-independent),  $A_t \sim \pi(a|s)$  for any  $t > 0$

# Policy in MDP

- 1 Given an MDP  $(S, A, P, R, \gamma)$  and a policy  $\pi$
- 2 The state sequence  $S_1, S_2, \dots$  is a Markov process  $(S, P^\pi)$
- 3 The state and reward sequence  $S_1, R_1, S_2, R_2, \dots$  is a Markov reward process  $(S, P^\pi, R^\pi, \gamma)$  where,

$$P^\pi(s'|s) = \sum_{a \in A} \pi(a|s) P(s'|s, a)$$
$$R^\pi(s) = \sum_{a \in A} \pi(a|s) R(s, a)$$

# Comparison of MP/MRP and MDP



# Value function for MDP

- ① The state-value function  $v_\pi(s)$  of an MDP is the expected return starting from state  $s$ , and following policy  $\pi$

$$v_\pi(s) = \mathbb{E}_\pi[G_t | s_t = s] \quad (3)$$

- ② The action-value function  $q_\pi(s, a)$  is the expected return starting from state  $s$ , taking action  $a$ , and then following policy  $\pi$

$$Q_\pi(s, a) = \mathbb{E}_\pi[G_t | s_t = s, A_t = a] \quad (4)$$

- ③ We have the relation between  $v_\pi(s)$  and  $Q_\pi(s, a)$

$$v_\pi(s) = \sum_{a \in A} \pi(a|s) Q_\pi(s, a) \quad (5)$$

# Bellman Expectation Equation

- 1 The state-value function can be decomposed into immediate reward plus discounted value of the successor state,

$$V_{\pi}(s) = E_{\pi}[R_{t+1} + \gamma V_{\pi}(s_{t+1}) | s_t = s] \quad (6)$$

- 2 The action-value function can similarly be decomposed

$$Q_{\pi}(s, a) = E_{\pi}[R_{t+1} + \gamma Q_{\pi}(s_{t+1}, A_{t+1}) | s_t = s, A_t = a] \quad (7)$$

# Bellman Expectation Equation for $V^\pi$ and $Q^\pi$

$$v_\pi(s) = \sum_{a \in A} \pi(a|s) q_\pi(s, a) \quad (8)$$

$$q_\pi(s, a) = R_s^a + \gamma \sum_{s' \in S} P(s'|s, a) v_\pi(s') \quad (9)$$


Thus

$$v_\pi(s) = \sum_{a \in A} \pi(a|s) (R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) v_\pi(s')) \quad (10)$$

$$q_\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \sum_{a' \in A} \pi(a'|s') q_\pi(s', a') \quad (11)$$




# Backup Diagram for $V^\pi$



$$v_\pi(s) = \sum_{a \in A} \pi(a|s) (R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) v_\pi(s')) \quad (12)$$

# Backup Diagram for $Q^\pi$




$$q_\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \sum_{a' \in A} \pi(a'|s') q_\pi(s', a') \quad (13)$$


# Policy Evaluation

- 1 Evaluate the value of state given a policy: compute  $v_{\pi}(s)$
- 2 Also called as prediction

# Example: Mars Rover




## Example: Policy Evaluation

$s_1$	$s_2$	$s_3$	$s_4$	$s_5$	$s_6$	$s_7$
						

- 1 Two actions: *Left* and *Right*
- 2 For all actions, reward: +5 in  $s_1$ , +10 in  $s_7$ , 0 in all other states. So that we can represent  $R = [5, 0, 0, 0, 0, 0, 10]$
- 3 Let's have a deterministic policy  $\pi(s) = \text{Left}$  and  $\gamma = 0$  for any state  $s$ , then what is the value of the policy?
- 4  $V^\pi = [5, 0, 0, 0, 0, 0, 10]$
- 5 Iterative:  $v_k^\pi(s) = r(s, \pi(s)) + \gamma \sum_{s' \in S} P(s'|s, \pi(s)) v_{k-1}^\pi(s')$

## Example: Policy Evaluation

$s_1$	$s_2$	$s_3$	$s_4$	$s_5$	$s_6$	$s_7$
						

- 1  $R = [5, 0, 0, 0, 0, 0, 10]$
- 2 Practice 1: Deterministic policy  $\pi(s) = \text{Left}$  with  $\gamma = 0.5$  for any state  $s$ , then what are the state values under the policy?
- 3 Practice 2: Stochastic policy  $P(\pi(s) = \text{Left}) = 0.5$  and  $P(\pi(s) = \text{Right}) = 0.5$  and  $\gamma = 0.5$  for any state  $s$ , then what are the state values under the policy?
- 4 Following the iteration:

$$v_k^\pi(s) = r(s, \pi(s)) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, \pi(s)) v_{k-1}^\pi(s')$$

# Decision Making in Markov Decision Process

- ① Prediction (evaluate a given policy):
  - ① Input: MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and policy  $\pi$  or MRP  $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$
  - ② Output: value function  $v_\pi$
- ② Control (search the optimal policy):
  - ① Input: MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
  - ② Output: optimal value function  $v_*$  and optimal policy  $\pi_*$
- ③ Prediction and control can be solved by dynamic programming.

# Dynamic Programming

Dynamic Programming is a very general solution method for problems which have two properties:

- ① Optimal substructure
  - ① Principle of optimality applies
  - ② Optimal solution can be decomposed into subproblems
- ② Overlapping subproblems
  - ① Subproblems recur many times
  - ② Solutions can be cached and reused

Markov decision processes satisfy both properties

- ① Bellman equation gives recursive decomposition
- ② Value function stores and reuses solutions



# Policy evaluation on MDP

- ① Problem: Evaluate a given policy  $\pi$  for a MDP
- ② Output the value function under policy  $v_\pi$
- ③ Solution: iteration on Bellman expectation backup
- ④ Synchronous backup algorithm:
  - ① At each iteration  $k+1$   
update  $v_{k+1}(s)$  from  $v_k(s')$  for all states  $s \in \mathcal{S}$  where  $s'$  is a successor state of  $s$

$$v_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) (R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) v_k(s')) \quad (14)$$

- ⑤ Convergence:  $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_\pi$

# Policy evaluation: Iteration on Bellman expectation backup

Bellman expectation backup for a particular policy


$$v_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) (R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) v_k(s')) \quad (15)$$

Or if in the form of MRP  $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}, \gamma \rangle$

$$v_{k+1}(s) = R^\pi(s) + \gamma P^\pi(s'|s) v_k(s') \quad (16)$$

# Evaluating a Random Policy in the Small Gridworld

Example 4.1 in the Sutton RL textbook.



- 1 Undiscounted episodic MDP ( $\gamma = 1$ )
- 2 Nonterminal states  $1, \dots, 14$
- 3 Two terminal states (two shaded squares)
- 4 Action leading out of grid leaves state unchanged,  $P(7|7, \text{right}) = 1$
- 5 Reward is  $-1$  until the terminal state is reached
- 6 Transition is deterministic given the action, e.g.,  $P(6|5, \text{right}) = 1$
- 7 Uniform random policy  $\pi(l|\cdot) = \pi(r|\cdot) = \pi(u|\cdot) = \pi(d|\cdot) = 0.25$

# Evaluating a Random Policy in the Small Gridworld

## 1 Iteratively evaluate the random policy

$V_k$  for the  
Random Policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

$k = 10$



0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

# Practice: Gridworld

## Textbook Example 3.5: GridWorld



3.3	8.8	4.4	5.3	1.5
1.5	3.0	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1.0	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2.0

# Optimal Value Function

- 1 The optimal state-value function  $v_*(s)$  is the maximum value function over all policies

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

- 2 The optimal policy

$$\pi_*(s) = \arg \max_{\pi} v_{\pi}(s)$$

- 3 An MDP is “solved” when we know the optimal value
- 4 There exists a unique optimal value function, but could be multiple optimal policies (two actions that have the same optimal value function)

# Finding Optimal Policy

- 1 An optimal policy can be found by maximizing over  $q_*(s, a)$ ,

$$\pi_*(a|s) = \begin{cases} 1, & \text{if } a = \arg \max_{a \in A} q_*(s, a) \\ 0, & \text{otherwise} \end{cases}$$

- 2 There is always a deterministic optimal policy for any MDP
- 3 If we know  $q_*(s, a)$ , we immediately have the optimal policy