

Lecture 6: Model-free Control

Bolei Zhou

The Chinese University of Hong Kong

bzhou@ie.cuhk.edu.hk

January 22, 2020

Today's Plan

① Last Time


- ① Model-free prediction: Estimate the value function of an **unknown** MDP
- ② Monte-Carlo (MC) and Temporal Difference (TD)

② This Time

- ① Model-free control: Optimize the value function of an **unknown** MDP
- ② Generalized Policy Iteration (GPI) with MC and TD


Review: Dynamic Programming Backup

$$v(S_t) \leftarrow \mathbb{E}_{\pi}[R_{t+1} + \gamma v(S_{t+1})]$$




Review: Monte-Carlo Backup

$$v(S_t) \leftarrow v(S_t) + \alpha(G_t - v(S_t))$$



Review: Temporal-Difference Backup

$$TD(0) : v(S_t) \leftarrow v(S_t) + \alpha(R_{t+1} + \gamma v(S_{t+1}) - v(S_t))$$



Control: Optimize the policy for a MDP



- ① Model-based control: optimize the value function with known MDP
- ② Model-free control: optimize the value function with unknown MDP
- ③ Many model-free RL examples: Go, robot locomotion, patient treatment, helicopter control, Atari, Starcraft

Policy Iteration

1 Iterate through the two steps:

- 1 **Evaluate** the policy π (computing v given current π)
- 2 **Improve** the policy by acting greedily with respect to v_π

$$\pi' = \text{greedy}(v_\pi) \quad (1)$$



Policy Iteration for a Known MDP

- 1 compute the state-action value of a policy π :

$$q_{\pi_i}(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) v_{\pi_i}(s')$$


- 2 Compute new policy π_{i+1} for all $s \in \mathcal{S}$ following

$$\pi_{i+1}(s) = \arg \max_a q_{\pi_i}(s, a) \quad (2)$$

- 3 Problem: What to do if there is neither $R(s, a)$ nor $P(s'|s, a)$ known/available?

Generalized Policy Iteration with Action-Value Function

Monte Carlo version of policy iteration



- 1 Policy evaluation: Monte-Carlo policy evaluation $Q = q_\pi$
- 2 Policy improvement: Greedy policy improvement?

$$\pi(s) = \arg \max_a q(s, a)$$

Monte Carlo with Exploring Starts

- 1 One assumption to obtain the guarantee of convergence in PI:
Episode has exploring starts
- 2 Exploring starts can ensure all actions are selected infinitely often

Monte Carlo ES (Exploring Starts), for estimating $\pi \approx \pi_*$

Initialize:

$\pi(s) \in \mathcal{A}(s)$ (arbitrarily), for all $s \in \mathcal{S}$

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Loop forever (for each episode):

Choose $S_0 \in \mathcal{S}, A_0 \in \mathcal{A}(S_0)$ randomly such that all pairs have probability > 0

Generate an episode from S_0, A_0 , following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:

Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$\pi(S_t) \leftarrow \text{argmax}_a Q(S_t, a)$

Monte Carlo with ϵ -Greedy Exploration

- ❶ Trade-off between exploration and exploitation (we will talk about this in later lecture)
- ❷ ϵ -Greedy Exploration: Ensuring continual exploration
 - ❶ All actions are tried with non-zero probability
 - ❷ With probability $1 - \epsilon$ choose the greedy action
 - ❸ With probability ϵ choose an action at random

$$\pi(a|s) = \begin{cases} \epsilon/|\mathcal{A}| + 1 - \epsilon & \text{if } a^* = \arg \max_{a \in \mathcal{A}} Q(s, a) \\ \epsilon/|\mathcal{A}| & \text{otherwise} \end{cases}$$

Monte Carlo with ϵ -Greedy Exploration

- ① Policy improvement theorem: For any ϵ -greedy policy π , the ϵ -greedy policy π' with respect to q_π is an improvement, $v_{\pi'}(s) \geq v_\pi(s)$

$$\begin{aligned} q_\pi(s, \pi'(s)) &= \sum_{a \in \mathcal{A}} \pi'(a|s) q_\pi(s, a) \\ &= \frac{\epsilon}{|\mathcal{A}|} \sum_{a \in \mathcal{A}} q_\pi(s, a) + (1 - \epsilon) \max_a q_\pi(s, a) \\ &\geq \frac{\epsilon}{|\mathcal{A}|} \sum_{a \in \mathcal{A}} q_\pi(s, a) + (1 - \epsilon) \sum_{a \in \mathcal{A}} \frac{\pi(a|s) - \frac{\epsilon}{|\mathcal{A}|}}{1 - \epsilon} q_\pi(s, a) \\ &= \sum_{a \in \mathcal{A}} \pi(a|s) q_\pi(s, a) = v_\pi(s) \end{aligned}$$

Therefore, $v_{\pi'}(s) \geq v_\pi(s)$ from the policy improvement theorem

Monte Carlo with ϵ -Greedy Exploration

Algorithm 1


```
1: Initialize  $Q(S, A) = 0, N(S, A) = 0, \epsilon = 1, k = 1$ 
2:  $\pi_k = \epsilon$ -greedy( $Q$ )
3: loop
4:   Sample  $k$ -th episode  $(S_1, A_1, R_2, \dots, S_T) \sim \pi_k$ 
5:   for each state  $S_t$  and action  $A_t$  in the episode do
6:      $N(S_t, A_t) \leftarrow N(S_t, A_t) + 1$ 
7:      $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{1}{N(S_t, A_t)}(G_t - Q(S_t, A_t))$ 
8:   end for
9:    $k \leftarrow k + 1, \epsilon \leftarrow 1/k$ 
10:   $\pi_k = \epsilon$ -greedy( $Q$ )
11: end loop
```

MC vs. TD for Prediction and Control

- ① Temporal-difference (TD) learning has several advantages over Monte-Carlo (MC)
 - ① Lower variance
 - ② Online
 - ③ Incomplete sequences
- ② So we can use TD instead of MC in our control loop
 - ① Apply TD to $Q(S, A)$
 - ② Use ϵ -greedy policy improvement
 - ③ Update every time-step rather than at the end of one episode

Recall: TD Prediction

- 1 An episode consists of an alternating sequence of states and state-action pairs:



- 2 TD(0) method for estimating the value function $V(S)$

$A_t \leftarrow$ action given by π for S


Take action A_t , observe R_{t+1} and S_{t+1}

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

- 3 How about estimating action value function $Q(S)$?

Sarsa: On-Policy TD Control

- 1 An episode consists of an alternating sequence of states and state-action pairs:



- 2 ϵ -greedy policy for one step, then bootstrap the action value function:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \right]$$

- 3 The update is done after every transition from a nonterminal state S_t
- 4 TD target $\delta_t = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1})$

Sarsa algorithm

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Repeat (for each step of episode):

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

n -step Sarsa

- 1 Consider the following n -step Q-returns for $n = 1, 2, \infty$

$$n = 1(\text{Sarsa}) q_t^{(1)} = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1})$$

$$n = 2 \quad q_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 Q(S_{t+2}, A_{t+2})$$

\vdots

$$n = \infty(\text{MC}) \quad q_t^\infty = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-t-1} R_T$$

- 2 Thus the n -step Q-return is defined as

$$q_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q(S_{t+n}, A_{t+n})$$

- 3 n -step Sarsa updates $Q(s,a)$ towards the n -step Q-return:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (q_t^{(n)} - Q(S_t, A_t))$$

On-policy Learning and Off-policy learning

- ① On-policy learning: Learn about policy π from the experience sampled from π
 - ① Behave non-optimally in order to explore all actions, then reduce the exploration. e.g., ϵ -greedy
- ② Another solution is to use two different policies:
 - ① one is learned about and becomes the optimal policy
 - ② the other one is more exploratory and is used to generate behavior
- ③ Off-policy learning: Learn about policy π from the experience sampled from another policy b
 - ① π : target policy
 - ② b : behavior policy

Off-policy Learning

- 1 Evaluate target policy $\pi(a|s)$ to compute $v_\pi(s)$ or $q_\pi(s, a)$, while following behaviour policy $\mu(a|s)$

$$S_1, A_1, R_2, \dots, S_T \sim \mu$$

Update π using $S_1, A_1, R_2, \dots, S_T$

- 2 Why is this important?
 - 1 Learn from observing humans or other agents
 - 2 Re-use experience generated from old policies $\pi_1, \pi_2, \dots, \pi_{t-1}$
 - 3 Learn about optimal policy while following exploratory policy

Importance Sampling

- 1 Estimate the expectation of a function

$$E_{x \sim P}[f(x)] = \int f(x)P(x)dx \approx \frac{1}{n} \sum_i f(x_i)$$

- 2 But sometimes it is difficult to sample x from $P(x)$, then we can sample x from another distribution $Q(x)$, then correct the weight

$$\begin{aligned}\mathbb{E}_{x \sim P}[f(x)] &= \int P(x)f(x)dx \\ &= \int Q(x)\frac{P(x)}{Q(x)}f(x)dx \\ &= \mathbb{E}_{x \sim Q}\left[\frac{P(x)}{Q(x)}f(x)\right] \approx \frac{1}{n} \sum_i \frac{P(x_i)}{Q(x_i)}f(x_i)\end{aligned}$$

Importance Sampling for Off-Policy RL

- 1 Estimate the expectation of return using trajectories sampled from another policy (behavior policy)

$$\begin{aligned}\mathbb{E}_{T \sim \pi}[g(T)] &= \int P(T)g(T)dT \\ &= \int Q(T)\frac{P(T)}{Q(T)}g(T)dT \\ &= \mathbb{E}_{T \sim \mu}\left[\frac{P(T)}{Q(T)}g(T)\right] \\ &\approx \frac{1}{n} \sum_i \frac{P(T_i)}{Q(T_i)}g(T_i)\end{aligned}$$

Importance Sampling for Off-Policy Monte Carlo

- 1 Generate episode from behavior policy μ and compute the generated return G_t

$$S_1, A_1, R_2, \dots, S_T \sim \mu$$

- 2 Weight return G_t according to similarity between policies
 - 1 Multiply importance sampling corrections along whole episode

$$G_t^{\pi/\mu} = \frac{\pi(A_t|S_t)}{\mu(A_t|S_t)} \frac{\pi(A_{t+1}|S_{t+1})}{\mu(A_{t+1}|S_{t+1})} \dots \frac{\pi(A_T|S_T)}{\mu(A_T|S_T)} G_t$$

- 3 Update value towards correct return

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t^{\pi/\mu} - V(S_t))$$

Importance Sampling for Off-Policy TD

- ① Use TD targets generated from μ to evaluate π
- ② Weight TD target $R + \lambda V(S')$ by importance sampling
- ③ Only need a single importance sampling correction

$$V(S_t) \leftarrow V(S_t) + \alpha \left(\frac{\pi(A_t|S_t)}{\mu(A_t|S_t)} (R_{t+1} + \lambda V(S_{t+1})) - V(S_t) \right)$$

- ④ Policies only need to be similar over a single step

Q Learning

- 1 Off-policy learning of action values $Q(s, a)$
- 2 No importance sampling is needed
- 3 Next action is chosen using behavior policy $A_{t+1} \sim \mu(.|S_t)$. However, we consider alternative action $A' \sim \pi(.|S_t)$
- 4 update $Q(S_t, A_t)$ towards value of alternative action

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A') - Q(S_t, A_t))$$

Off-Policy Control with Q-Learning

- 1 We allow both behavior and target policies to improve
- 2 The target policy π is **greedy** on $Q(s, a)$

$$\pi(S_{t+1}) = \arg \max_{a'} Q(S_{t+1}, a')$$

- 3 The behavior policy μ is **ϵ -greedy** on $Q(s, a)$
- 4 Thus Q-learning target:

$$\begin{aligned} R_{t+1} + \gamma Q(S_{t+1}, A') &= R_{t+1} + \gamma Q(S_{t+1}, \arg \max_{a'} Q(S_{t+1}, a')) \\ &= R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a') \end{aligned}$$

- 5 Thus the Q-Learning update

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

Q-learning algorithm

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Repeat (for each step of episode):

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

Comparison of Sarsa and Q-Learning

1 Sarsa: On-Policy TD control

Choose action A_t from S_t using policy derived from Q with ϵ -greedy

Take action A_t , observe R_{t+1} and S_{t+1}

Choose action A_{t+1} from S_{t+1} using policy derived from Q with ϵ -greedy

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

2 Q-Learning: Off-Policy TD control

Choose action A_t from S_t using policy derived from Q with ϵ -greedy



Take action A_t , observe R_{t+1} and S_{t+1}

Take action A_{t+1} from S_{t+1} using policy derived from Q with greedy

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

Comparison of Sarsa and Q-Learning

1 Backup diagram for Sarsa and Q-learning




Summary of DP and TD

Expected Update (DP)	Sample Update (TD)
Iterative Policy Evaluation $V(s) \leftarrow \mathbb{E}[R + \gamma V(S') s]$	TD Learning $V(S) \leftarrow^\alpha R + \gamma V(S')$
Q-Policy Iteration $Q(S, A) \leftarrow \mathbb{E}[R + \gamma Q(S', A') s, a]$	Sarsa $Q(S, A) \leftarrow^\alpha R + \gamma Q(S', A')$
Q-Value Iteration $Q(S, A) \leftarrow \mathbb{E}[R + \gamma \max_{a' \in \mathcal{A}} Q(S', A') s, a]$	Q-Learning $Q(S, A) \leftarrow^\alpha R + \gamma \max_{a' \in \mathcal{A}} Q(S', a')$

where $x \leftarrow^\alpha y$ is defined as $x \leftarrow x + \alpha(y - x)$

Unified View of Reinforcement Learning



Sarsa and Q-Learning Example

`https:
//github.com/cuhkrlcourse/RLexample/tree/master/modelfree`

To Do for the CNY

- ① Reinforce yourself with Chapter 1 to Chapter 8 (done)
 - ① We will get into Part II: Approximate Solution Methods after CNY
- ② Finish your Assignment 1